



# **FAULT TOLERANCE – RELIABLE GROUP COMMUNICATION**

**timofei.istomin@unitn.it**

**Thanks for the slides to:**


**Giuliano Mega**

**Davide Frey**

**Luca Mottola**

# RELIABLE GROUP COMMUNICATION

- It is of critical importance if we want to exploit process resilience by replication
- Classical example: the *state machine* approach;
  - make your process a *deterministic state machine*, so that...
  - ... given  $n$  copies of your process, if you feed them the same inputs:
    1. they will emit the same outputs
    2. their internal state will be the same



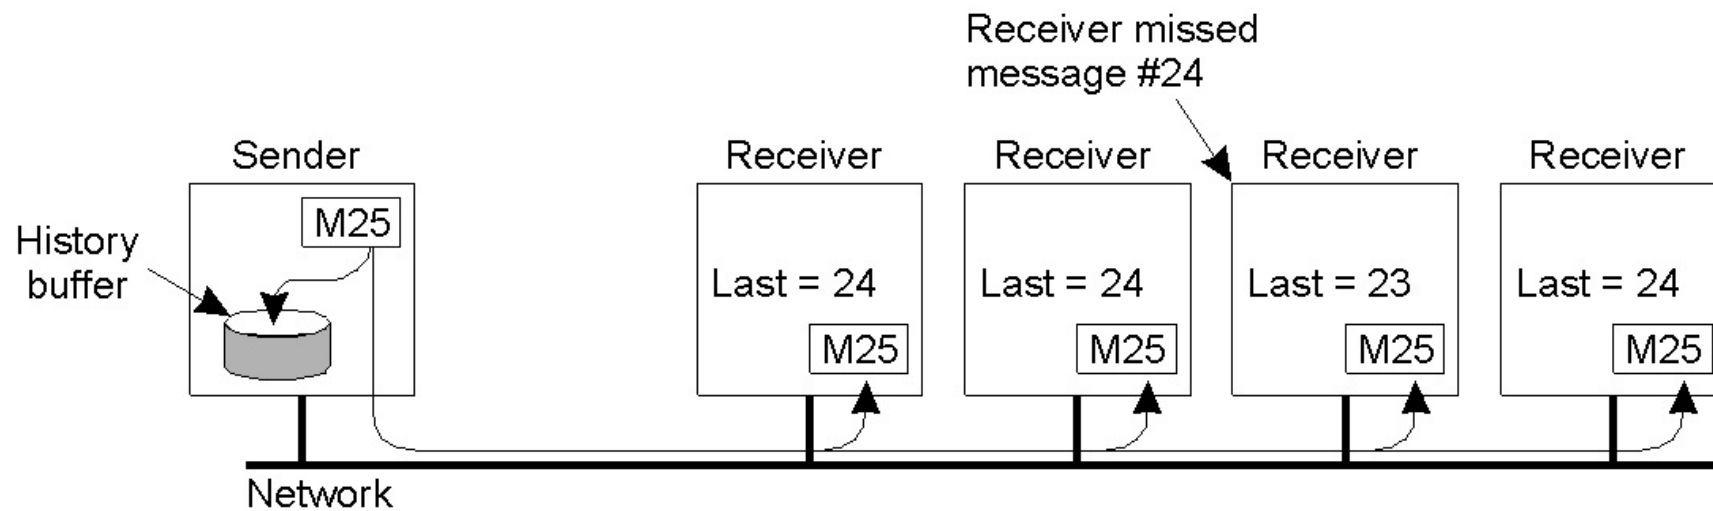
# RELIABLE GROUP COMMUNICATION

- Reliable multicast is surprisingly tricky
  - Yet, we need it to deliver messages to all members in a process group
- Usually, all we have is reliable point-to-point service or unreliable multicast service
  - Achieving reliable multicast through multiple reliable point-to-point channels may be not efficient
- Plus...
  - ... what happens if the sender crashes while sending?
  - What if the group changes while a message is being sent? Who should get the message, and who should not?

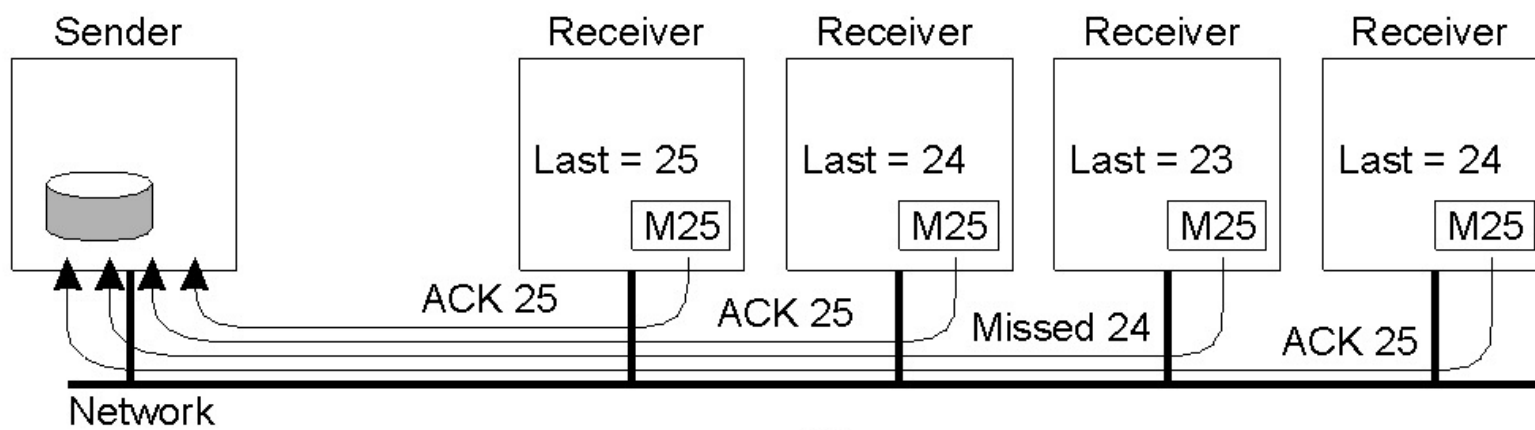
# RELIABLE GROUP COMMUNICATION

- We look at two cases for reliable multicast
- **Case 1:** groups are fixed (i.e., processes do not crash), communication failures only
  - All group members should receive the multicast
  - Unreliable multicast primitive is available
- **Case 2:** faulty processes, but reliable channels
- We start from **Case 1**

# BASIC RELIABLE MULTICAST



(a)



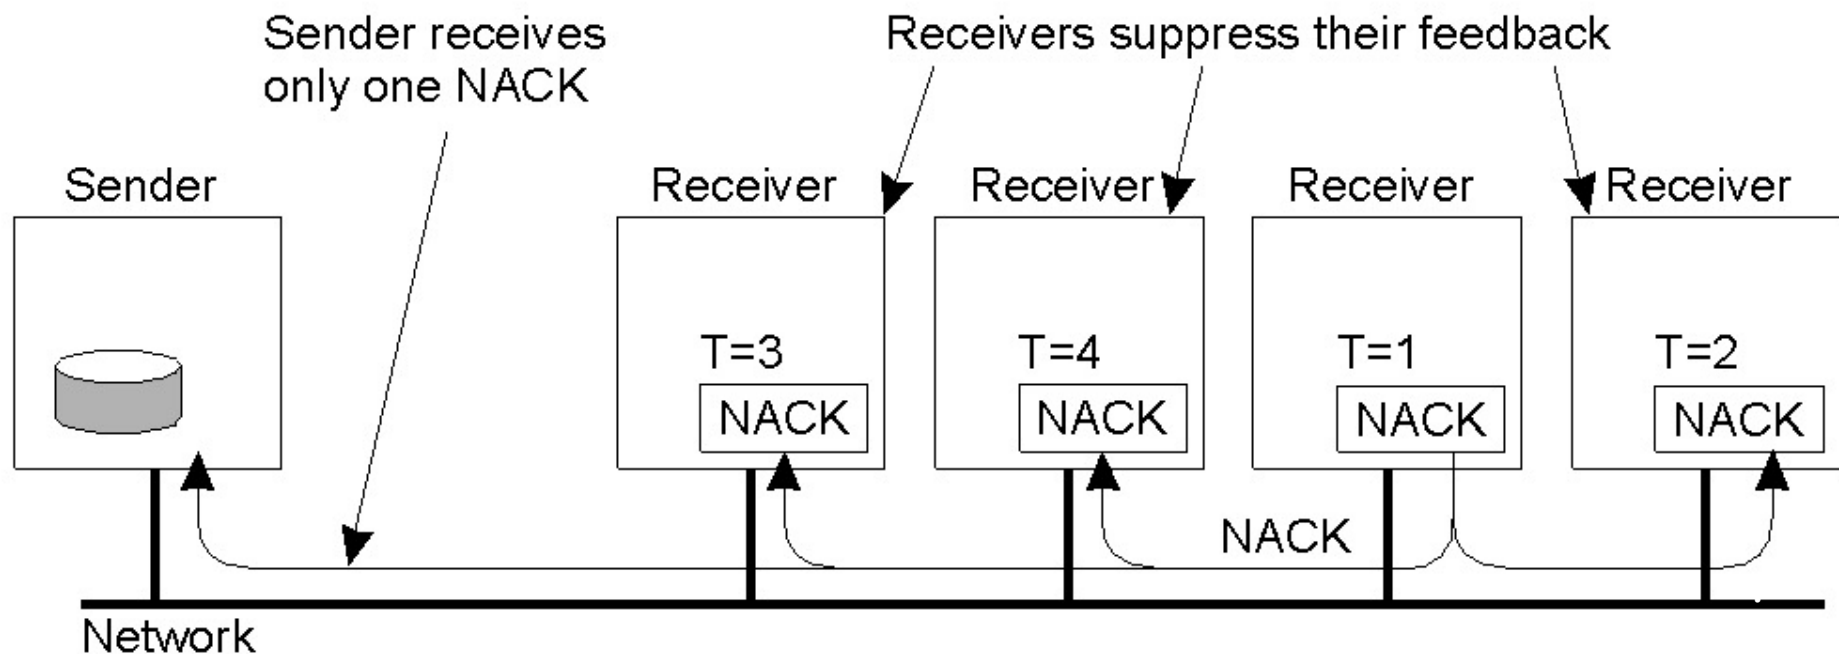
(b)

# BASIC RELIABLE MULTICAST

- What if we make the receivers ACK all messages?
- Key problem: too many ACKs:
  - “feedback implosion”
- Mitigation strategies:
  - piggyback ACK messages on traffic (only works if traffic is constant enough), and
  - retransmit using point-to-point primitive
- Scalability problems still persist
- Negative ACKs: only send feedback if message not received
  - sender might have to cache messages forever
  - still may result in too many ACKs

# SCALABLE RELIABLE MULTICAST

- A first solution is non-hierarchical feedback control, implemented in SRM
  - Negative acknowledgements are multicast
  - Randomly staggered
- Used in practice in several applications




# SCALABLE RELIABLE MULTICAST

- Staggering “the right amount” is non-trivial
  - depends on network delays, right value might be group-dependent, especially over WANs
- Processes that don't lose messages still receive NACKs from those who do
  - Dynamically create separate groups for processes that tend to miss messages together, difficult to achieve in large-scale settings
- Ultimately, flat groups will always have scalability limitations



# HIERARCHICAL FEEDBACK CONTROL

- With Hierarchical Feedback control, receivers are organized in groups headed by a coordinator



# HIERARCHICAL FEEDBACK CONTROL


- The coordinator can adopt any strategy within its group
- In addition it can request retransmissions to its parent coordinator
  - A coordinator can remove a message from its buffer if it has received an ACK from
    - All the receivers in its group
    - All of its child coordinators
- The problem is that the hierarchy has to be constructed and maintained
  - Difficult to implement on physical networks, usually done via application-level overlays

# RELIABLE MULTICAST

- Recall – two cases:
  - **Case 1:** fixed groups, unreliable channels;
  - **Case 2:** dynamic groups (processes can join, leave or crash), reliable channels
- Now we look at **Case 2**
- The main problem is **dynamic membership**:
  - group membership may change while multicasts are being issued;
  - to simplify the application programmer's life the middleware can provide **virtual synchrony (VS)**


# VIEWS AND VIEW SYNCHRONY

- In VS, each process maintains a “view” of the group
  - i.e., local knowledge of who is part of the group, **and therefore to whom messages should be delivered to**
- Views “get replaced” as the group changes: processes *install new views* when a new member joins the group or someone leaves the group



# VIEWS AND VIEW SYNCHRONY

- The system guarantees (under some assumptions) **that all correct processes will see the same sequence of group views**
- Therefore, each view defines a global **epoch**




- Another important property guaranteed by virtual synchrony: **multicasts do not cross epoch boundaries**



# VIEWS AND VIEW SYNCHRONY

- Why crossing epoch boundaries is a problem?




R does not have the same information as P does when it gets informed about Q's crash

P and R are not sure what information S have received since it has joined

# CLOSE SYNCHRONY


- Virtual synchrony *emulates* close synchrony where the events (multicasts, group changes) are *instantaneous*
- Easy to reason about, since crashes or joins never overlap with message delivery
- Multicasts also reliable: always get to everyone in the group





# VIEWS AND VIEW SYNCHRONY



- “Multicasts do not cross epoch boundaries”
  - If there are multicasts for messages  $m_1$  and  $m_2$ , and a view change occurs ...
  - ... to satisfy view synchrony, these multicasts **must complete** before the new view is installed



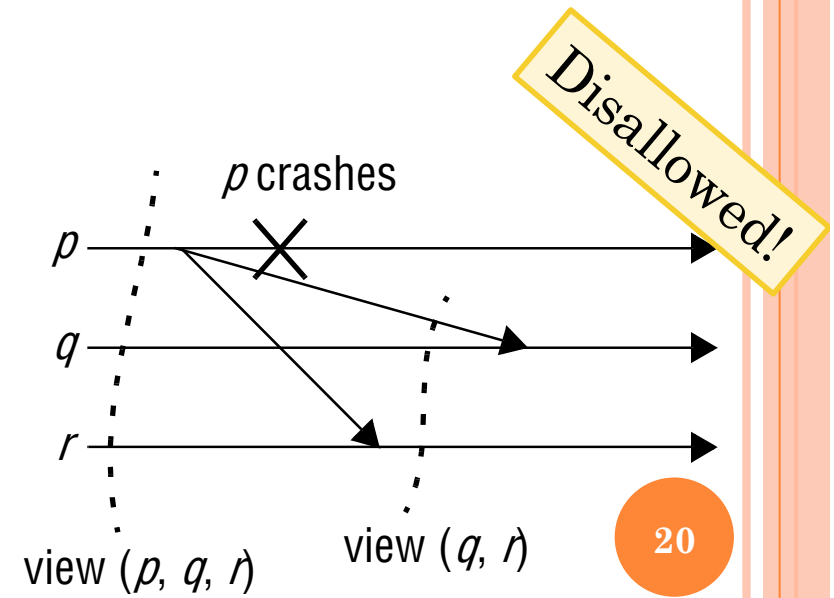
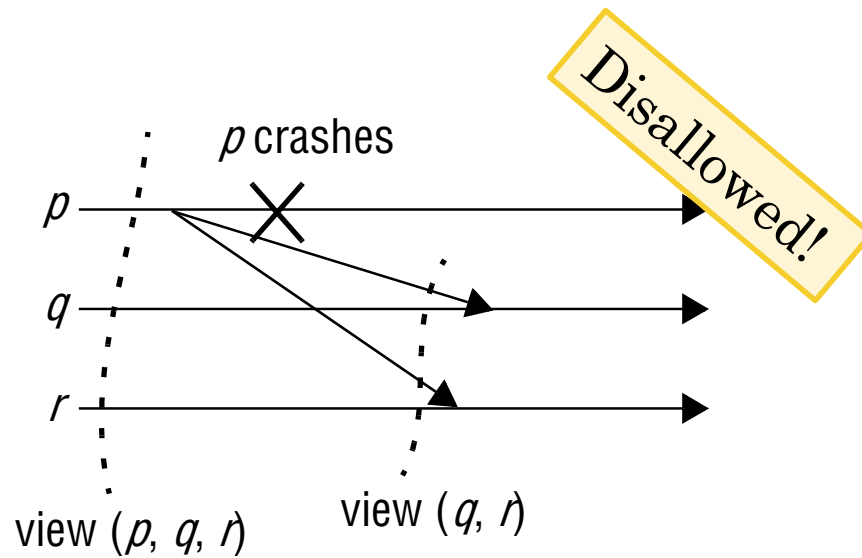


# MULTICASTS IN VIRTUAL SYNCHRONY

- What does it mean for multicast to “complete”?  
Who has to get it or not get it?
- Virtual synchrony multicasts are **reliable**
- A message multicast in an epoch  $E$  defined by view  $V$  should either:
  - be **delivered** within  $E$  by all **operational** participants
  - be **delivered** by none of the **operational** participants
  - One of these conditions is met: multicast is complete
- To exclude trivial protocols (just discard everything), the “*or none*” case is only allowed if the sender crashes
  - All “evidence” erased: as if it had never been sent
  - We only care about **operational** participants...

# EXAMPLES



# EXAMPLES



# THE FLAVORS OF VIRTUAL SYNCHRONY

- In virtual synchrony, messages can't cross epochs ...
- ... but what can happen inside an epoch?
- Three base flavors:
  - unordered;
  - FIFO ordered;
  - causally ordered;
- where each of those can be:
  - totally ordered;
  - not totally ordered;
- for a total of six flavors

**Detour: slides about  
message ordering...**



# NAMING THE FLAVORS

- The six versions of virtually synchronous reliable multicast

<b>Multicast</b>	<b>Basic Message Ordering</b>	<b>Total-ordered Delivery?</b>
Reliable multicast	None	No
FIFO multicast	FIFO-ordered delivery	No
Causal multicast	Causal-ordered delivery	No
Atomic multicast	None	Yes
FIFO atomic multicast	FIFO-ordered delivery	Yes
Causal atomic multicast	Causal-ordered delivery	Yes

# EXERCISE 1

Is this execution virtually synchronous? What about FIFO, causally, or total ordered? Assume all start sharing the same view V0.

## P1

1. delivers M1
2. installs view V1
3. delivers M2
4. multicasts M3
5. delivers M3
6. installs view V2

## P2

1. multicasts M1
2. delivers M1
3. installs view V1
4. multicasts M2
5. delivers M2
6. delivers M3
7. installs view V2

## P3

1. delivers M1
2. installs view V1
3. delivers M3
4. delivers M2
5. installs view V2

## P4

1. installs view V2
2. delivers M2


### Group views

V0={P1,P2,P3,P4}

V1={P1,P2,P3}

V2={P1,P2,P3,P4}

## EXERCISE 1



- **not** virtually synchronous
- FIFO
- **not** causal
- **not** total



## EXERCISE 2

Is this execution virtually synchronous? What about FIFO, causally, or total ordered? Assume all start sharing the same view V0.

### P1

1. multicasts M1
2. delivers M1
3. installs view V1
4. delivers M2
5. installs view V2
6. delivers M3

### P2

1. delivers M1
2. installs view V1
3. delivers M2
4. installs view V2
5. delivers M3

### P3

1. installs view V2
2. multicasts M3
3. delivers M2
4. delivers M3

### P4

1. delivers M1
2. installs view V1
3. multicasts M2
4. delivers M2
5. installs view V2
6. delivers M3


### Group views

V0={P1,P2,P3,P4}

V1={P1,P2,P4}

V2={P1,P2,P3,P4}

## EXERCISE 2



- **not** virtually synchronous
- FIFO
- causal
- total

# IMPLEMENTING VIRTUAL SYNCHRONY


- We have seen the execution model...
- ... but how can you enforce it in a real system?
  - ISIS toolkit, now called **VSync** [Birman et. al 1991]
  - JGroups (part of JBoss)
- **Assumptions:**
  - Reliable FIFO channels
  - Processes may silently crash

# MULTICAST STABILITY

- Multicast is implemented as a sequence of unicasts to all group members
- Channels are reliable, therefore all the unicasts eventually reach their destination
  - ... unless the destination crashes
- The only case in which a multicast can fail is when the initiator crashes in the middle of it
- To overcome this case, every process  $p$ 
  - Delivers immediately the message  $m$  it received, but ...
  - .. keeps a copy of  $m$  until it is **sure** that every *correct* process in the *current view* has it ...
  - ... once this happens,  $m$  is said to be ***stable***, and the copy can be dropped
  - sends the copy to processes that might need it


# DETERMINING MESSAGE STABILITY

- Based on assumptions, the initiator learns that its multicast  $m$  is **stable** after it has sent the individual (unicast) messages to all the group members



# ANNOUNCING MESSAGE STABILITY

- To inform the other nodes (and free their memory by deleting their copy of  $m$ ), the initiator announces to the group that  $m$  is now stable
  - Piggybacking this information on the following multicasts or sending a dedicated one, if needed



## HANDLING GROUP CHANGES

- When a process leaves, joins, or detects a crash, it sends a view change message to the group
- Upon receiving a view change with view  $V_1$ :
  - until the new view is installed:
    - pause sending *new* multicasts
    - keep delivering incoming messages sent in the old view, but
    - defer messages sent in the new view
  - do an “all-to-all” echo: send *all* **unstable** messages to all processes in the **new** view  $V_1$ , followed by a **FLUSH** message
  - wait until the **FLUSH** message arrives from **every** other process in  $V_1$
  - install the view  $V_1$

Unstable messages are ignored by a joining node; only the **FLUSH** messages matter to it


## HANDLING GROUP CHANGES

- Having received **FLUSH** from all processes in  $V_1$ , we know that we have received all the unstable messages sent in  $V_0$  from all currently operational processes (because of FIFO), therefore
  - we deliver them in the old view  $V_0$
  - we install the new view  $V_1$
- If someone else crashes before sending **FLUSH**: repeat the all-to-all phase with the next view  $V_2$




# EXAMPLE EXECUTION

initiator crashes, processes  
1, 2, 3 deliver  $m$  in  $V_0$



crash detected in  $V_0$



- ➡
- all receive  $V_1$
  - all send to all:  
 $m$ ,  $\text{FLUSH}(V_0)$




- 4 and 5 deliver  $m$  in  $V_0$
- all wait till  $\text{FLUSH}(V_0)$   
comes from all in  $V_1$
- all install  $V_1$

$V_0 = \{0, 1, 2, 3, 4, 5\}$

$V_1 = \{1, 2, 3, 4, 5\}$

# EXAMPLE EXECUTION


crash detected in  $V_0$



What if another node crashes during all-to-all phase?

- all receive  $V_1$ ,
  - all send to all:  $m$ ,  $\text{FLUSH}(V_0)$
- ➡
- 4 delivers  $m$
  - All wait till  $\text{FLUSH}(V_0)$  comes from **all** in  $V_1$ ... bad luck!

Another crash detected in  $V_0$



➡ all send to all:  $V_2$ ,  $\text{FLUSH}(V_1)$  ➡


- all receive  $V_2$
- wait till  $\text{FLUSH}(V_1)$  comes from all in  $V_2$
- all install  $V_1$
- all install  $V_2$

$V_1 = \{1, 2, 3, 4, 5\}$

$V_2 = \{1, 2, 3, 4\}$

## A NOTE ON MULTICAST STABILITY

- Note that the described algorithm delivers arriving messages right away, without ensuring that they are stable, therefore, the following is possible




- The VS properties still hold (see previous slides)
- However, if this behaviour is undesirable, it can be solved by adding a second phase to each multicast...



## A TWO-PHASE VARIANT

- The initiator collects acknowledgements about each multicast and confirms the stability by sending a second-phase message
- The receivers deliver the message only after they have learnt that the message is stable



This will slow down the protocol significantly!

# LITERATURE

## Reliable group communication:

- G. Coulouris, J. Dollimore, T. Kindberg. Distributed Systems: Concepts and Design (5<sup>th</sup> edition). Addison-Wesley, 2012
- A. S. Tanenbaum, M. van Steen. Distributed Systems: Principles and Paradigms (2<sup>nd</sup> edition)
- <http://jgroups.org/manual4/index.html>